

Lab Eleven – Point Estimation and Resampling

Tips

- Complete the tasks below. Make sure to start your solutions in on a new line that starts with “**Solution:**”.
- Make sure to use the Quarto Cheatsheet. This will make completing and writing up the lab *much* easier.
- Make sure to use “enter” to make your code readable, so it doesn’t run off the page. I switched the Quarto document to automatically render to .pdf, so you can see what you’re submitting by default. You can switch back for the highlighting by moving the html block above the pdf block in the YAML header.
- Don’t forget that you can copy-and-paste code when you’re doing something similar as we did in class!
- Make sure to plot all computed probabilities. This is good `ggplot` practice AND will help make sure you don’t make a mistake when computing the probability.

1 Introduction

When performing point estimation, there isn’t just one *right* answer. For example, the Method of Moments (MOM) and Maximum Likelihood Estimates (MLEs) aren’t always the same.

One question that comes up, is why we calculate the sample variance as

$$S_{n-1}^2 = \frac{1}{n-1} \sum_{i=1}^n (X_i - \bar{X})^2,$$

when it is supposed to be the “average squared distance from the mean, i.e.,

$$S_n^2 = \frac{1}{n} \sum_{i=1}^n (X_i - \bar{X})^2.$$

2 Comparing Estimators

2.1 Initial Simulation

In this lab, we will explore why we use S_{n-1}^2 and not S_n^2 . To do this, we will generate data from four known distributions:

- Poisson($\lambda = 2$)
- Binomial($n = 50, p = \frac{5-\sqrt{21}}{10}$)
- Normal($\mu = 0, \sigma = \sqrt{2}$)
- Exponential($\lambda = \frac{1}{\sqrt{2}}$).

Remark: Note that I have selected the parameters so that the true variance for all four models is 2.

For each distribution, write a loop that completes the following 1000 times. Put `set.seed(7272)` at the top of your code chunk.

- Generate a sample of size n from the distribution.
- Calculate and save S_{n-1}^2 and S_n^2 .

Start with $n = 20$. Is there evidence that S_{n-1}^2 is a better estimator than S_n^2 ?

Create a plot or combination of plots (using patchwork) to show the results across all four distributions. Make sure to create a corresponding table that numerically summarizes the results of the simulation. You should include

$$\text{Bias} = \text{Average of Estimates} - \text{True Value}$$

$$\text{Precision} = \frac{1}{\text{Variance of Estimates}}$$

$$\text{Mean Square Error} = \text{Bias}^2 + \text{Variance of Estimates}.$$

Solution

```

1 set.seed(7272)
2
3 #numbers
4 n = 20
5 true_val = 2
6
7 pois_n = rep(NA, 1000)
8 pois_n1 = rep(NA, 1000)
9
10 for (i in 1:1000){
11   x_pois = rpois(n, 2)
12   pois_n1[i] = var(x_pois)
13   pois_n[i] = (n-1)/n * var(x_pois)
14 }
15
16 binom_n = rep(NA, 1000)
17 binom_n1 = rep(NA, 1000)
18
19 for (i in 1:1000){
20   x_binom = rbinom(n, size = 50, prob = (5-sqrt(21))/10)
21   binom_n1[i] = var(x_binom)
22   binom_n[i] = (n-1)/n * var(x_binom)
23 }
24
25 norm_n = rep(NA, 1000)
26 norm_n1 = rep(NA, 1000)
27
28 for (i in 1:1000){
29   x_norm = rnorm(n, mean = 0, sd = sqrt(2))
30   norm_n1[i] = var(x_norm)
31   norm_n[i] = (n-1)/n * var(x_norm)
32 }
33
34 expo_n = rep(NA, 1000)
35 expo_n1 = rep(NA, 1000)
36
37 for (i in 1:1000){
38   x_expo = rexp(n, rate = sqrt(1/2)) #dis is the same as 1/sqrt(2) but im cool like that
39   expo_n1[i] = var(x_expo)
40   expo_n[i] = (n-1)/n * var(x_expo)
41 }
42
43 results = bind_rows(
44   tibble(distribution = "Poisson", estimator = "S2_n-1", estimate = pois_n1),
45   tibble(distribution = "Poisson", estimator = "S2_n", estimate = pois_n),
46   tibble(distribution = "Binomial", estimator = "S2_n-1", estimate = binom_n1),
47   tibble(distribution = "Binomial", estimator = "S2_n", estimate = binom_n),


```

```

48 tibble(distribution = "Normal", estimator = "S2_n-1", estimate = norm_n1),
49 tibble(distribution = "Normal", estimator = "S2_n", estimate = norm_n),
50 tibble(distribution = "Exponential", estimator = "S2_n-1", estimate = expo_n1),
51 tibble(distribution = "Exponential", estimator = "S2_n", estimate = expo_n)
52 )
53
54 ggplot(results, aes(x = estimator, y = estimate)) +
55   geom_boxplot() +
56   facet_wrap(~distribution, scales = "free")+
57   geom_hline(yintercept=true_val)+
58   xlab("Estimator") +
59   ylab("Estimated Variance")+
60   theme_bw()

```

Figure 1: Variance Estimators by Distribution



```

1 summary_table = results |>
2   group_by(distribution, estimator) |>
3   summarize(
4     bias = mean(estimate) - true_val,
5     precision = 1/(var(estimate)),
6     mse = bias^2 + var(estimate),
7     average_estimate = mean(estimate), #added these bc relevant to simulations
8     variance_of_estimates = var(estimate),
9   )
10
11 knitr::kable(summary_table, digits = 4)

```

distribution	estimator	bias	precision	mse	average_estimate	variance_of_estimates
Binomial	S2_n	-0.1336	2.3821	0.4376	1.8664	0.4198
Binomial	S2_n-1	-0.0354	2.1498	0.4664	1.9646	0.4652
Exponential	S2_n	-0.1150	0.7746	1.3042	1.8850	1.2910
Exponential	S2_n-1	-0.0158	0.6991	1.4307	1.9842	1.4305

distribution	estimator	bias	precision	mse	average_estimate	variance_of_estimates
Normal	S2_n	-0.0978	2.2480	0.4544	1.9022	0.4448
Normal	S2_n-1	0.0023	2.0289	0.4929	2.0023	0.4929
Poisson	S2_n	-0.1153	2.3283	0.4428	1.8847	0.4295
Poisson	S2_n-1	-0.0161	2.1013	0.4762	1.9839	0.4759

Based on our above figure, we can see that for most distributions S^2_{n-1} is generally closer to the true variance value compared to the other estimator, S^2_n . We can generally keep an eye on this as evidence for our best estimator as of yet.

2.2 Simulation over various sample sizes

Consider the effect of the sample size have on these results. Provide a graphical and numerical summary of what happens as n increases – say $n = 20$, $n = 50$, $n = 100$, and $n = 500$. Does the result in your initial simulation hold across samples? Or, do things change as the sample size changes? Ensure to set your randomization seed.

Solution

```

1 set.seed(7272)
2
3 simulate_by_size = function(n){
4
5     pois_n = rep(NA, 1000)
6     pois_n1 = rep(NA, 1000)
7
8     for (i in 1:1000){
9         x_pois = rpois(n, 2)
10        pois_n1[i] = var(x_pois)
11        pois_n[i] = (n-1)/n * var(x_pois)
12    }
13
14    binom_n = rep(NA, 1000)
15    binom_n1 = rep(NA, 1000)
16
17    for (i in 1:1000){
18        x_binom = rbinom(n, size = 50, prob = (5-sqrt(21))/10)
19        binom_n1[i] = var(x_binom)
20        binom_n[i] = (n-1)/n * var(x_binom)
21    }
22
23    norm_n = rep(NA, 1000)
24    norm_n1 = rep(NA, 1000)
25
26    for (i in 1:1000){
27        x_norm = rnorm(n, mean = 0, sd = sqrt(2))
28        norm_n1[i] = var(x_norm)
29        norm_n[i] = (n-1)/n * var(x_norm)
30    }
31
32    expo_n = rep(NA, 1000)
33    expo_n1 = rep(NA, 1000)
34
35    for (i in 1:1000){
36        x_expo = rexp(n, rate = sqrt(1/2)) #dis is the same as 1/sqrt(2) but im cool like that
37        expo_n1[i] = var(x_expo)
38        expo_n[i] = (n-1)/n * var(x_expo)
39    }

```

```

40
41 results = bind_rows(
42   tibble(sample_size = n,distribution ="Poisson", estimator ="S2_n-1", estimate =pois_n1),
43   tibble(sample_size = n,distribution ="Poisson", estimator ="S2_n", estimate =pois_n),
44   tibble(sample_size = n,distribution ="Binomial", estimator ="S2_n-1", estimate =binom_n1),
45   tibble(sample_size = n,distribution ="Binomial", estimator ="S2_n", estimate =binom_n),
46   tibble(sample_size = n,distribution ="Normal", estimator ="S2_n-1", estimate =norm_n1),
47   tibble(sample_size = n,distribution ="Normal", estimator ="S2_n", estimate =norm_n),
48   tibble(sample_size = n,distribution ="Exponential", estimator ="S2_n-1", estimate =expo_n1),
49   tibble(sample_size = n,distribution ="Exponential", estimator ="S2_n", estimate =expo_n)
50 )
51 return(results)
52 }
53
54 all_results = bind_rows(
55   simulate_by_size(20),
56   simulate_by_size(50),
57   simulate_by_size(100),
58   simulate_by_size(500)
59 )
60
61 summary_table = all_results |>
62   group_by(sample_size, distribution, estimator) |>
63   summarize(
64     bias = mean(estimate) - true_val,
65     precision = 1 / var(estimate),
66     mse = bias^2 + var(estimate),
67     average_estimate = mean(estimate),
68     variance_of_estimates = var(estimate)
69   )
70
71 knitr::kable(summary_table, digits = 4)

```

sample_size	distribution	estimator	bias	precision	mse	average_estimate	variance_of_estimates
20	Binomial	S2_n	-0.1336	2.3821	0.4376	1.8664	0.4198
20	Binomial	S2_n-1	-0.0354	2.1498	0.4664	1.9646	0.4652
20	Exponential	S2_n	-0.1150	0.7746	1.3042	1.8850	1.2910
20	Exponential	S2_n-1	-0.0158	0.6991	1.4307	1.9842	1.4305
20	Normal	S2_n	-0.0978	2.2480	0.4544	1.9022	0.4448
20	Normal	S2_n-1	0.0023	2.0289	0.4929	2.0023	0.4929
20	Poisson	S2_n	-0.1153	2.3283	0.4428	1.8847	0.4295
20	Poisson	S2_n-1	-0.0161	2.1013	0.4762	1.9839	0.4759
50	Binomial	S2_n	-0.0374	5.2558	0.1917	1.9626	0.1903
50	Binomial	S2_n-1	0.0026	5.0477	0.1981	2.0026	0.1981
50	Exponential	S2_n	0.0087	1.4557	0.6870	2.0087	0.6869
50	Exponential	S2_n-1	0.0497	1.3981	0.7177	2.0497	0.7153
50	Normal	S2_n	-0.0304	6.2500	0.1609	1.9696	0.1600
50	Normal	S2_n-1	0.0098	6.0025	0.1667	2.0098	0.1666
50	Poisson	S2_n	-0.0293	4.6500	0.2159	1.9707	0.2151
50	Poisson	S2_n-1	0.0109	4.4659	0.2240	2.0109	0.2239
100	Binomial	S2_n	-0.0476	10.9444	0.0936	1.9524	0.0914
100	Binomial	S2_n-1	-0.0279	10.7266	0.0940	1.9721	0.0932
100	Exponential	S2_n	0.0076	2.9136	0.3433	2.0076	0.3432
100	Exponential	S2_n-1	0.0278	2.8556	0.3510	2.0278	0.3502
100	Normal	S2_n	-0.0045	12.4465	0.0804	1.9955	0.0803


sample_size	distribution	estimator	bias	precision	mse	average_estimate	variance_of_estimates
100	Normal	S2_n-1	0.0157	12.1988	0.0822	2.0157	0.0820
100	Poisson	S2_n	-0.0110	10.1671	0.0985	1.9890	0.0984
100	Poisson	S2_n-1	0.0091	9.9648	0.1004	2.0091	0.1004
500	Binomial	S2_n	-0.0001	54.6630	0.0183	1.9999	0.0183
500	Binomial	S2_n-1	0.0039	54.4446	0.0184	2.0039	0.0184
500	Exponential	S2_n	-0.0042	15.5185	0.0645	1.9958	0.0644
500	Exponential	S2_n-1	-0.0002	15.4565	0.0647	1.9998	0.0647
500	Normal	S2_n	0.0005	58.6278	0.0171	2.0005	0.0171
500	Normal	S2_n-1	0.0045	58.3935	0.0171	2.0045	0.0171
500	Poisson	S2_n	-0.0025	54.1056	0.0185	1.9975	0.0185
500	Poisson	S2_n-1	0.0015	53.8894	0.0186	2.0015	0.0186

```

1 bias_plot = ggplot(summary_table,
2                     aes(x = sample_size, y = bias, color = estimator)) +
3   geom_line() +
4   geom_point() +
5   facet_wrap(~distribution) +
6   xlab("Sample Size") +
7   ylab("Bias")+
8   geom_hline(yintercept=0)+
9   scale_x_continuous(breaks = c(20, 50, 100, 500)) +
10  theme_bw()+
11  theme( #had to look this up so it wasnt in the way and the x intercept was visible
12        legend.position = "bottom",
13        axis.text.x = element_text(angle = 55, hjust = 1)
14  )
15
16 precision_plot = ggplot(summary_table,
17                          aes(x = sample_size, y = precision, color = estimator)) +
18  geom_line() +
19  geom_point() +
20  facet_wrap(~distribution) +
21  xlab("Sample Size")+
22  ylab("Precision")+
23  scale_x_continuous(breaks = c(20, 50, 100, 500)) +
24  theme_bw() +
25  theme(
26    legend.position = "bottom",
27    axis.text.x = element_text(angle = 55, hjust = 1)
28  )
29
30 bias_plot | precision_plot

```

Figure 2: Sample Size vs Bias and Precision



Looking first at our precision graphs, we see that both estimators evolve almost exactly the same when scaling with sample size, even having extremely similar values themselves for our precision measurement. Thus, we disregard the precision graph and use the bias graph to make a proper distinction between each estimator. We see that generally S^2_{n-1} has lower bias compared to S^2_n at every low sample size, thus we can say that it is a better estimator at low sample sizes. At large sample sizes, both estimators seem to converge very close to 0 bias, thus using either estimator at large sample sizes may not actually matter much.


2.3 Simulation

You have written code to draw a random sample of various sizes from each distribution, and to calculate and store the calculated estimates. Rework your code to produce a density plot of the s^2_{n-1} and s^2_n . Do you have a guess about their distribution? Ensure to set your randomization seed.

Solution

```
1 set.seed(7272)
2
3 ggplot(all_results, aes(x = estimate, fill = estimator)) +
4   geom_density(alpha = 0.4) +
5   facet_grid(sample_size~distribution, scales = "free") + #goated oscar refine
6   geom_vline(xintercept = true_val, linetype = "dashed") +
7   xlab("Estimated Variance") +
8   ylab("Density") +
9   theme_bw() +
10  theme(
11    legend.position = "bottom"
12  )
```

Figure 3: Density Plots of Variance Estimators with Varying Sample Sizes



Here, we see that at lower values of n , S^2_{n-1} is more centered around the true variance value. We also see that at higher n values, both estimators share the same shape and are pretty gaussian distributed around the true value. I could have guessed that at lower n values, the density plot was going to share similarities with a chi squared plot as we are sampling on variance, which implies a chi squared distribution with right skewness. This is exactly what we see here.

2.4 Bootstrap

Now, take *one* sample from each distribution. Instead of drawing a random sample at each iteration, sample *with replacement* from the one sample. Produce a density plot of the s^2_{n-1} and s^2_n . Do we get roughly the same information as we do in the full simulation? Try this over several seeds – what do you notice? Ensure to set your randomization seed.

Solution

```
1 set.seed(7272)
2
3 boot_by_size = function(n){
4
5 #distributions
6   x_pois = rpois(n, 2)
```



```

7   x_binom = rbinom(n, size = 50, prob = (5-sqrt(21))/10)
8   x_norm = rnorm(n, mean = 0, sd = sqrt(2))
9   x_expo = rexp(n, rate = sqrt(1/2))
10
11  #storage
12  pois_n = rep(NA, 1000)
13  pois_n1 = rep(NA, 1000)
14
15  #booting here
16  for (i in 1:1000){
17    boot_pois = sample(x_pois, size = n, replace = TRUE)
18    pois_n1[i] = var(boot_pois)
19    pois_n[i] = (n-1)/n * var(boot_pois)
20  }
21
22  #rinse repeat for each distr
23  binom_n = rep(NA, 1000)
24  binom_n1 = rep(NA, 1000)
25
26  for (i in 1:1000){
27    boot_binom = sample(x_binom, size = n, replace = TRUE)
28    binom_n1[i] = var(boot_binom)
29    binom_n[i] = (n-1)/n * var(boot_binom)
30  }
31
32  norm_n = rep(NA, 1000)
33  norm_n1 = rep(NA, 1000)
34
35  for (i in 1:1000){
36    boot_norm = sample(x_norm, size = n, replace = TRUE)
37    norm_n1[i] = var(boot_norm)
38    norm_n[i] = (n-1)/n * var(boot_norm)
39  }
40
41  expo_n = rep(NA, 1000)
42  expo_n1 = rep(NA, 1000)
43
44  for (i in 1:1000){
45    boot_expo = sample(x_expo, size = n, replace = TRUE)
46    expo_n1[i] = var(boot_expo)
47    expo_n[i] = (n-1)/n * var(boot_expo)
48  }
49
50  results = bind_rows(
51    tibble(sample_size = n, distribution = "Poisson", estimator = "S2_n-1", estimate = pois_n1),
52    tibble(sample_size = n, distribution = "Poisson", estimator = "S2_n", estimate = pois_n),
53    tibble(sample_size = n, distribution = "Binomial", estimator = "S2_n-1", estimate = binom_n1),
54    tibble(sample_size = n, distribution = "Binomial", estimator = "S2_n", estimate = binom_n),
55    tibble(sample_size = n, distribution = "Normal", estimator = "S2_n-1", estimate = norm_n1),
56    tibble(sample_size = n, distribution = "Normal", estimator = "S2_n", estimate = norm_n),
57    tibble(sample_size = n, distribution = "Exponential", estimator = "S2_n-1", estimate = expo_n1),
58    tibble(sample_size = n, distribution = "Exponential", estimator = "S2_n", estimate = expo_n)
59  )
60
61  return(results)
62 }


```

```

63 bootstrap_results = bind_rows(
64   boot_by_size(20),
65   boot_by_size(50),
66   boot_by_size(100),
67   boot_by_size(500)
68 )
69
70 ggplot(bootstrap_results, aes(x = estimate, fill = estimator)) +
71   geom_density(alpha = 0.4) +
72   facet_grid(sample_size ~ distribution, scales = "free") +
73   geom_vline(xintercept = true_val, linetype = "dashed") +
74   xlab("Estimated Variance") +
75   ylab("Density") +
76   theme_bw() +
77   theme(
78     legend.position = "bottom"
79   )
80

```

Figure 4: Bootstrap Density Plots of Variance Estimators by Sample Size (Seed: 7272)




```

1 set.seed(5055)
2
3 bootstrap_results = bind_rows(
4   boot_by_size(20),
5   boot_by_size(50),
6   boot_by_size(100),
7   boot_by_size(500)
8 )
9
10 ggplot(bootstrap_results, aes(x = estimate, fill = estimator)) +
11   geom_density(alpha = 0.4) +
12   facet_grid(sample_size ~ distribution, scales = "free") +
13   geom_vline(xintercept = true_val, linetype = "dashed") +
14   xlab("Estimated Variance") +
15   ylab("Density") +
16   theme_bw() +
17   theme(
18     legend.position = "bottom"
19   )

```

Figure 5: Bootstrap Density Plots of Variance Estimators by Sample Size (Seed: 5055)



3 Executive Summary

Summarize the work you've done here to provide a brief (200-300 word) summary of the simulation and the results using your work in Lab 11. Your summary should be professional, clear, and all claims should be corroborated with statistical support.